

Progress Report (Jul 1 - Jul 8, 2020)

Doing

Done

- Compared GHOST against other baselines
 - beats DODGE, as in [Excel sheet](#)
 - loses against [Lin Tan's results](#) (wins in 2/10 datasets on recall)
 - Q: unsure if I should run GHOST for precision and compare, same for F1 score
- Ran modified version of SVM: [results](#) | [code](#) | [theory](#)
 - When tuned, beats DODGE in recall, has only one hyperparameter to tune (therefore tuning is fast), but achieves good precision in 5/10 datasets, and matches DODGE in popt20 in all
 - advantage over DODGE: one single learner with one hyperparameter means runtime << DODGE
 - Also applied kernel: surprisingly, linear kernel works better than rbf; possibly because:
 - SE data is low dimensional, but the rbf kernel projects to an infinite-dimensional space
 - the k hyperparameter is lost in the projection and therefore is now redundant
- Implemented tab domination for [DANCE](#)
- [Baselines for issue close time](#) (AUC): ran DL and DODGE
 - Different people use different metrics
 - DODGE uses d2h and popt20
 - FFT uses accuracy, precision, recall, d2h, f1-score

Todo

- Run [FFT baselines](#) on issue close time
- Compare results between DL, DODGE, FFT
- Get DODGE pf results

Roadblocks

- Did not understand the data or results in [Kikas paper](#); [their code](#) is mostly Bash and some Java + Weka